AEP FROM ZERO • PART 4 • MULTIPLY THE TEAM (AI)

EPISODE 35

AI Agents and MCP Explained for Marketing Teams

AI agents explained for marketing teams: what is the difference between an AI assistant and an AI agent, and what is MCP, the Model Context Protocol, in simple words? Bloom's team runs a campaign-brief agent, traces its tool calls in a run log, adds a 48-hour frequency guardrail tool and tests it in dry-run mode, where the guardrail trips and changes the brief. You will learn tool calling and AI agent guardrails: narrow permissions, blast radius, stopping conditions, human-in-the-loop approval

Free course on iTV • Surge Software Solutions, Bengaluru • Bloom is fictional

Previously: an assistant that only talks

- Last time: RAG assistant, at least 8 of 10, cited
- Fixed by re-chunking: retrieval, not the model
- But it only answers; it never does the work
- Today: an agent drafts the Monday campaign brief
- And the brake: permissions, limits, approvals, logs

KEY IDEA

Bloom's RAG assistant answers honestly, but it never acts. Today an AI does the Monday legwork, and the team builds the brake first.

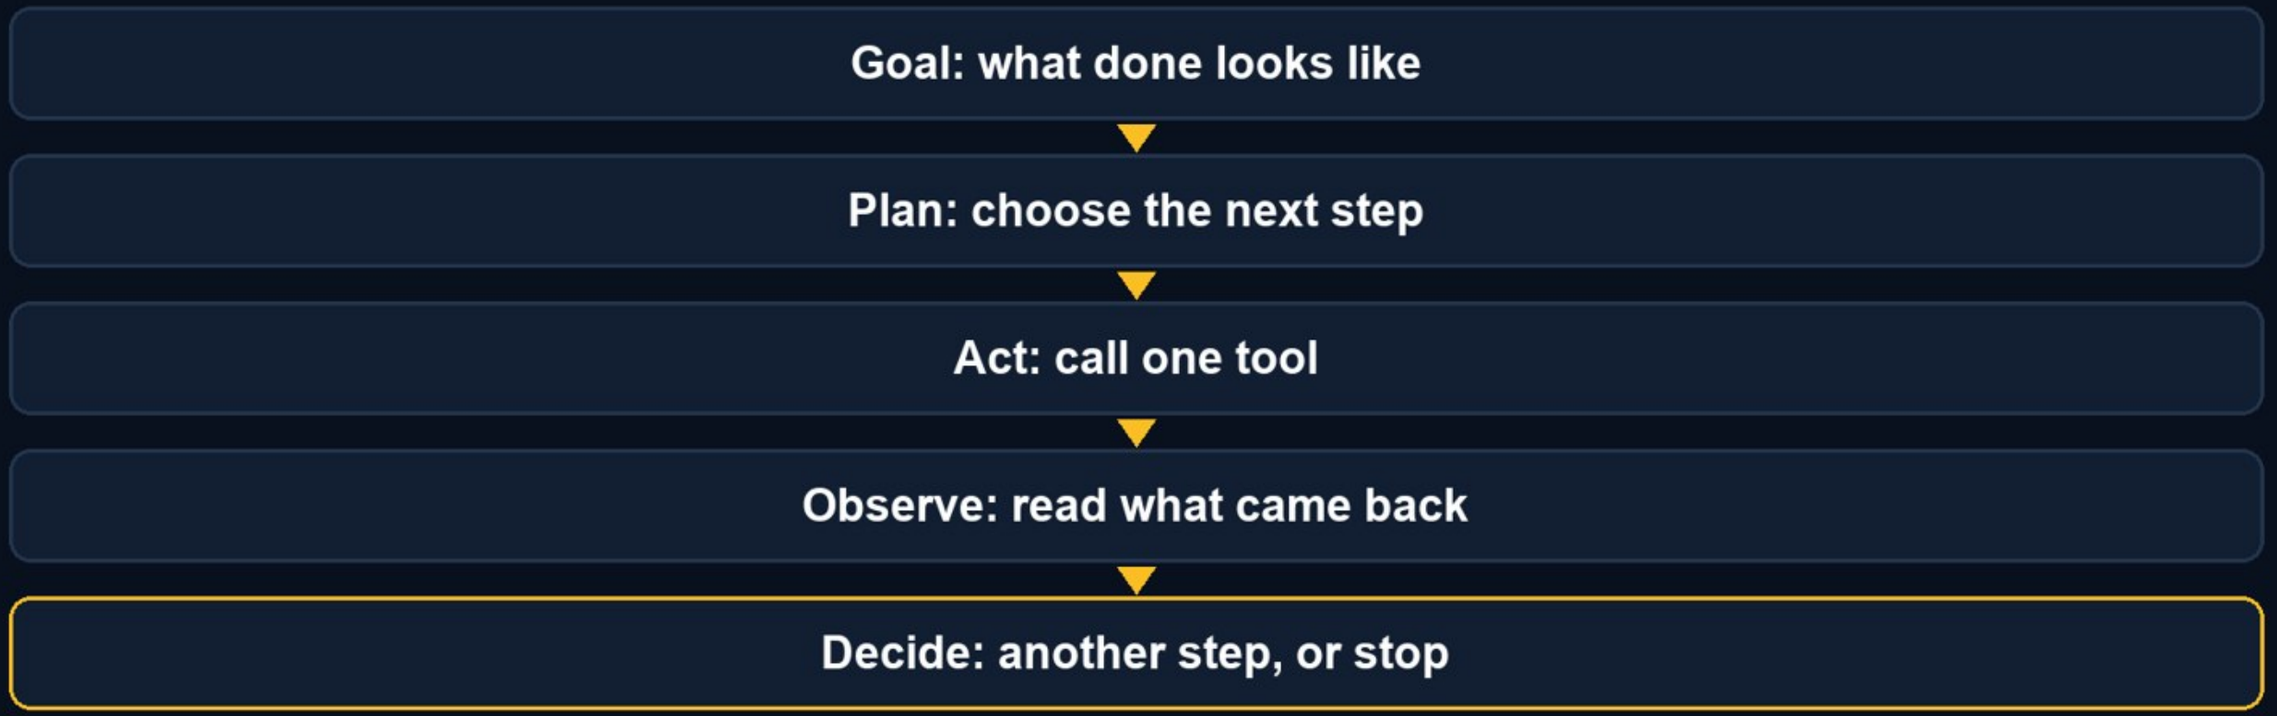
Assistant vs agent: what is the difference?

AI assistant	AI agent
Answers when you ask	Works towards a goal you set
One question, one reply	Many steps: plan, act, observe
You do the work afterwards	It calls tools to do the work
Risk: a wrong answer	Risk: a wrong action
Value: faster thinking	Value: hours of legwork saved

KEY IDEA

An assistant answers when asked. An agent plans, acts with tools, observes the result and continues until the goal is met or a limit is hit.

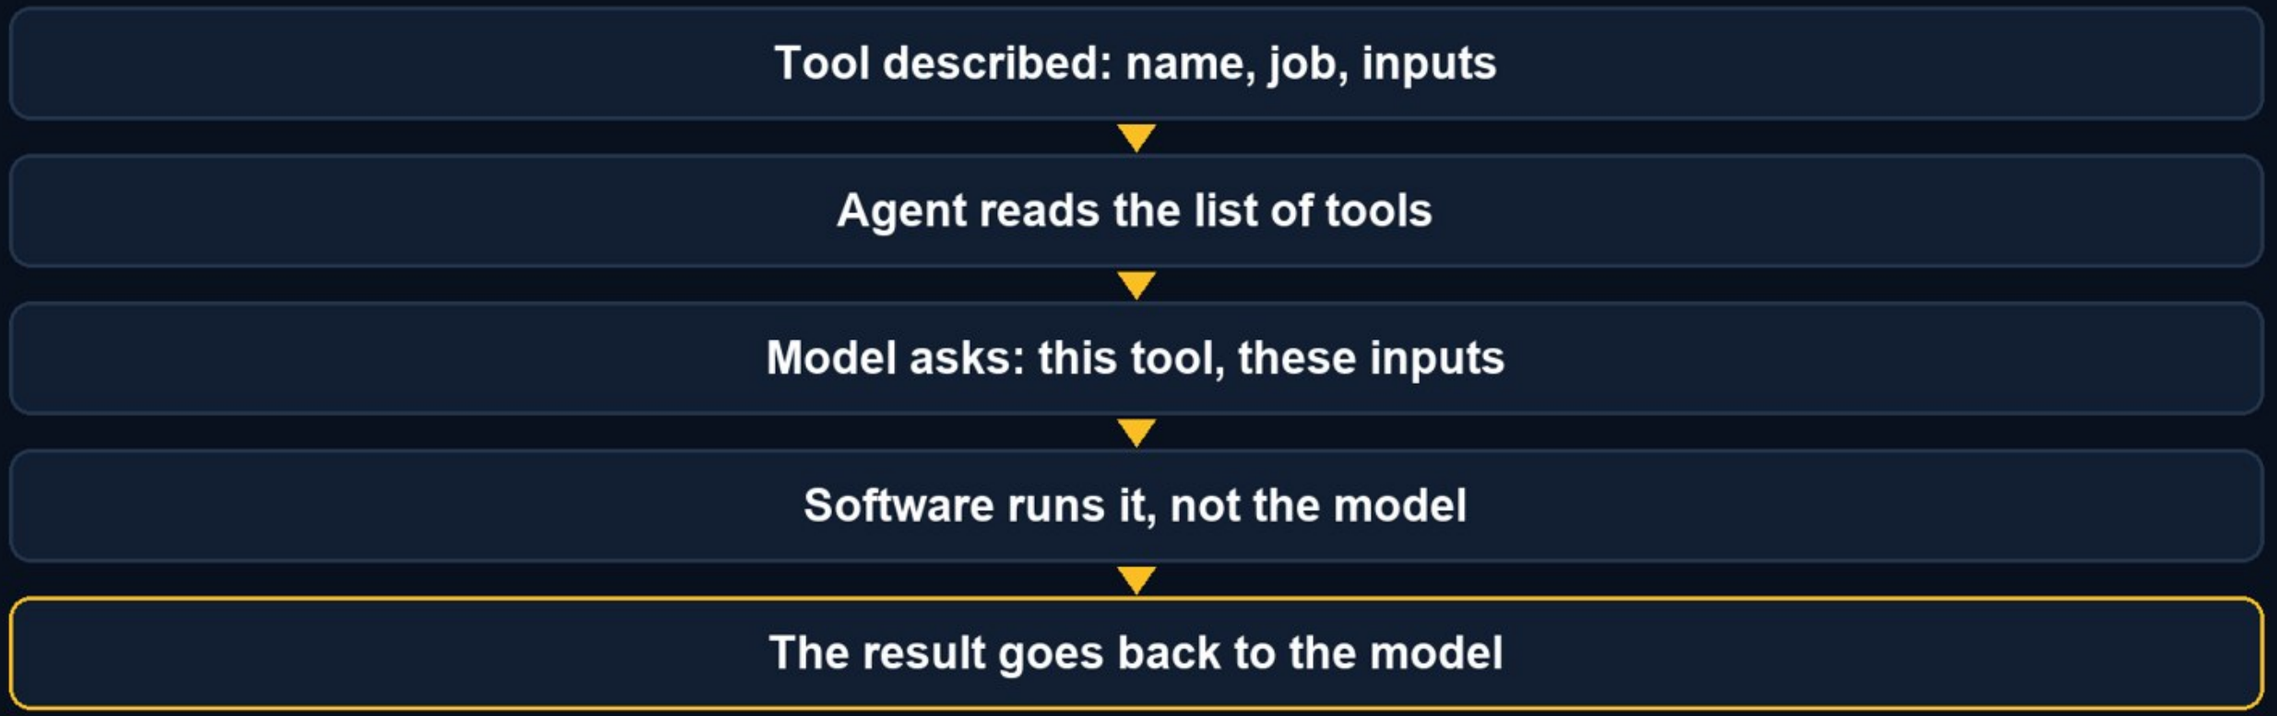
The agent loop and a definition of done



KEY IDEA

An agent runs a loop: plan, call a tool, read the result, decide the next step. It stops only when done, or when a limit is reached.

What is MCP? Tools and tool calling



KEY IDEA

A tool is a capability an agent may call. MCP is a standard way to describe tools, so any agent that speaks it can discover and call them.

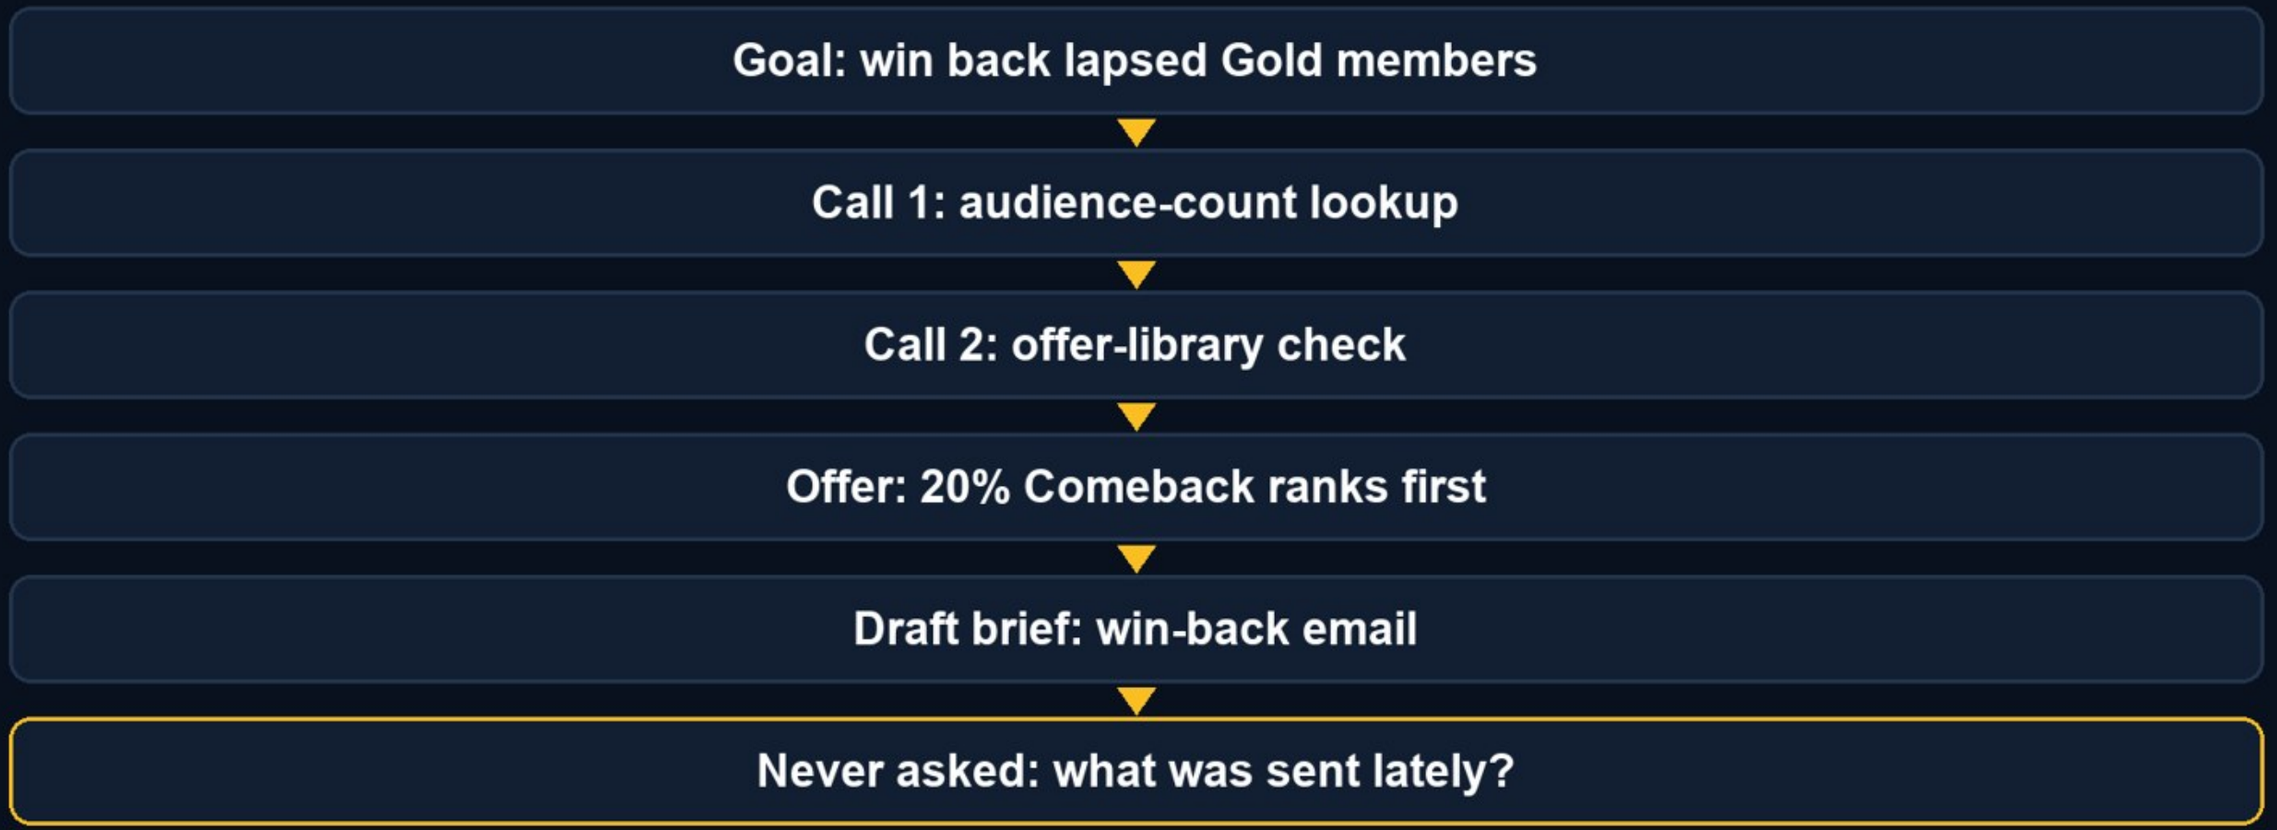
Bloom's campaign-brief agent and its tools

Tool	What it does	Access
Audience-count lookup	Counts people in an audience	Read-only
Offer-library check	Lists offers they qualify for	Read-only
Brief writer	Saves the draft brief	Write, one folder

KEY IDEA

Bloom's agent runs in a notebook outside the Adobe screens. It has three tools: an audience-count lookup, an offer-library check and a brief writer.

Reading the run log



KEY IDEA

A run log lists every tool call with its inputs and outputs. Read it to see what the agent did, and what it never checked.

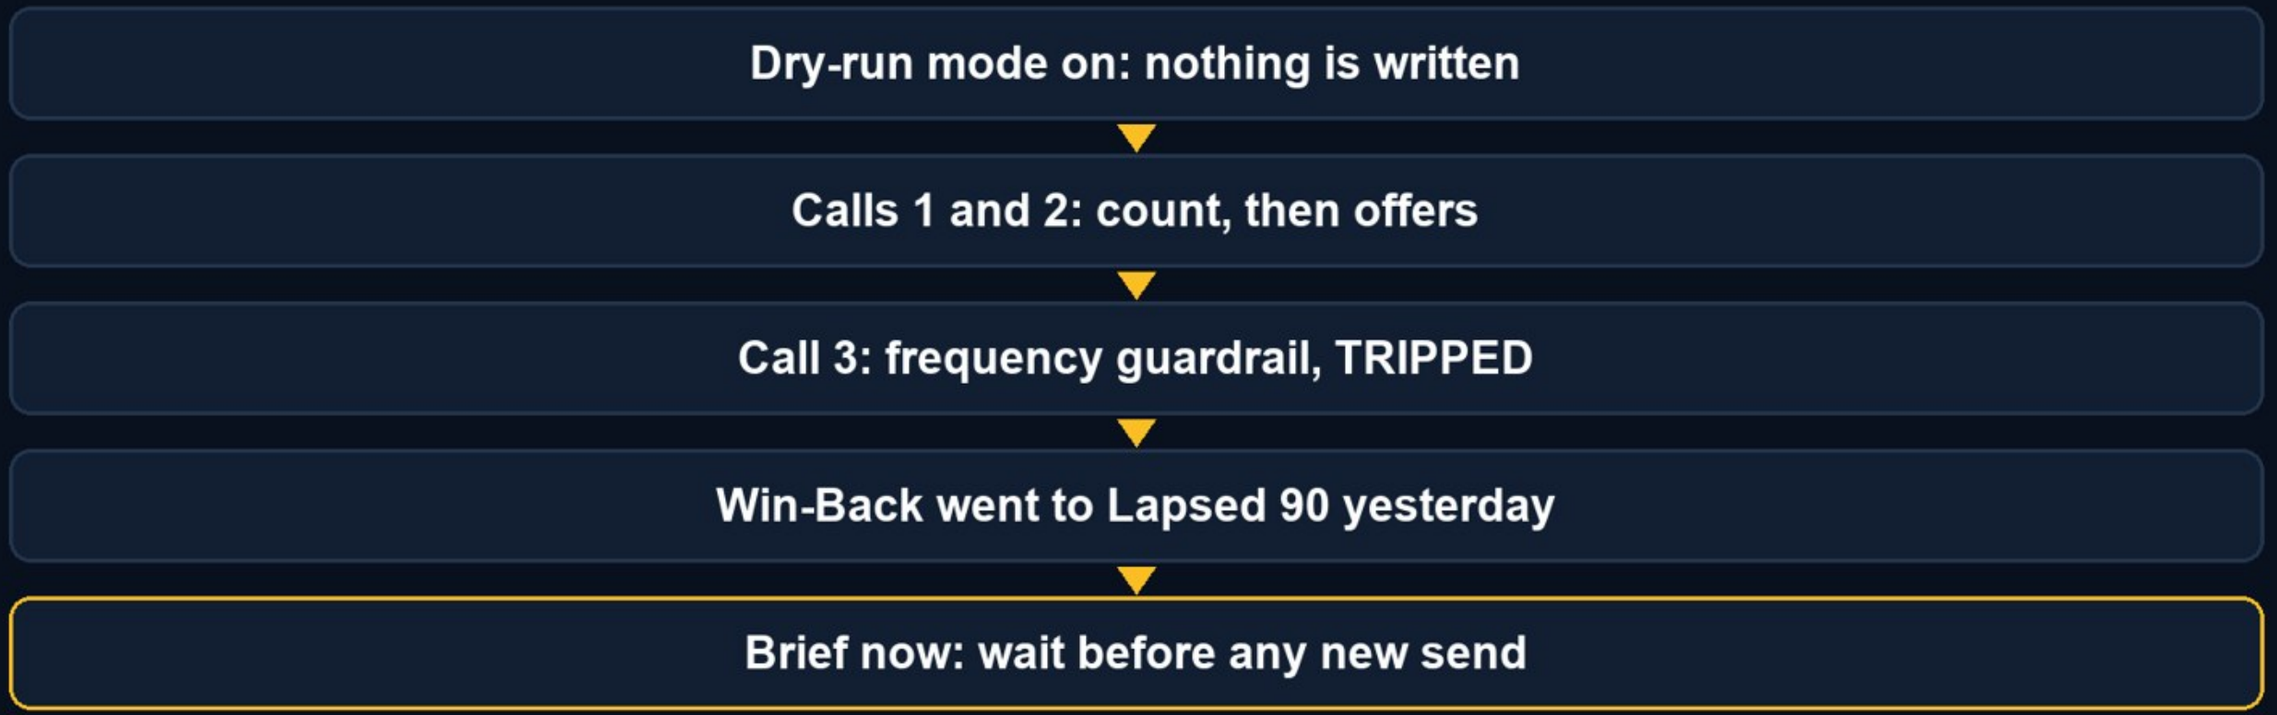
Divya's rule becomes a tool

Tool contract	Frequency-guardrail check
Input	Audience name, a 48-hour window
Check	Bloom's send log for that audience
Output	Clear, or a warning with the last send
Effect on the draft	On a warning: recommend waiting

KEY IDEA

Divya's rule becomes a frequency-guardrail tool: it reads the send log and warns if the audience was messaged in the last 48 hours.

Dry run: the guardrail trips



KEY IDEA

In dry-run mode the agent plans and reports but changes nothing. The guardrail tripped, because Bloom Win-Back went to Lapsed 90 yesterday.

How do you make an AI agent safe at work?

- Narrowest permissions that still finish the job
- Ask: what is the worst it could do if it misreads?
- Dry-run mode first: plan and report, change nothing
- Stopping conditions: steps, spend, time, done
- Human approval before anything irreversible
- Log every action, its inputs and its reason

KEY IDEA

Build the brake before the engine: narrow permissions, a blast-radius question, dry runs, stopping conditions, approval gates and a log of everything.

Narrow permissions and the blast radius

Tool	Access	Worst case if wrong
Audience-count lookup	Read-only	A wrong count in a draft
Offer-library check	Read-only	A wrong offer suggested
Frequency guardrail	Read-only	A missed warning
Brief writer	Write, one folder	A bad draft, thrown away

KEY IDEA

Give each tool the narrowest access that still finishes the job, then ask what the worst case would be if the agent misunderstood.

Stopping conditions: stop and report

- Maximum steps: say, for example, ten tool calls
- Maximum spend on model calls per run
- A timeout for the whole run
- Done: audience, offer, recommendation, checks logged
- On failure: stop and report, never retry forever

KEY IDEA

Every agent needs limits: maximum steps, maximum spend, a timeout and a definition of done. When it fails, it should stop and report.

Approval gates without approval fatigue

Step	Human approval?
Look up audience counts	No: read-only, and logged
Check the offer library	No: read-only, and logged
Write the draft brief	No: a draft is reversible
Send, delete or export	Yes: irreversible, Divya decides

KEY IDEA

Put the human where the action is irreversible: before a send, a delete or an export. Gating every step trains people to approve without reading.

Audit trails and the governance note

- What the agent may do alone
- What needs a human approval, and whose
- Dry-run mode as the default
- Stopping conditions and failure behaviour
- The log: every call, input, output and reason
- Who reads the log, and how often

KEY IDEA

Log what the agent did, why and with what inputs. Then write one page: what it may do alone, what needs approval, its limits and its audit.

Interview corner

- Put the human approval where it is irreversible
- Assistant answers; an agent plans, acts, observes
- Story: a 48-hour guardrail tripped in a dry run
- Brake: permissions, limits, one gate, a full log

KEY IDEA

Put the human approval where the action is irreversible. Then prove it with one story: a guardrail that tripped in a dry run.

Recap, your question, and Demo Day

- Assistant answers; agent plans, acts, observes
- MCP describes tools; you still grant the access
- Guardrail tripped in dry run: wait before sending
- Brake: permissions, limits, approval, logs
- Next: Demo Day, one customer, seven minutes

KEY IDEA

An agent acts, so build the brake first. Next: the day before Demo Day, and the whole build walked end to end.